

DATA GOVERNANCE AND SECURITY STANDARDS

Document ID: POL-TECH-302

Owner: Vikram Malhotra (CTO)

Classification: Internal Restricted

API schemas must follow OpenAPI specifications. Every service must support health check endpoints, structured JSON logging, and distributed tracing metadata headers (X-Request-ID).

In this context, the microservices architecture employs Kong API Gateway for routing client calls, securing endpoints via OAuth 2.0, and applying rate limits to protect backend services from denial of service.

Consequently, data pipelines ingest real-time sales transactions, process them via stream processors, and sync them to PostgreSQL and ChromaDB databases for reporting and analytics.

In this context, API schemas must follow OpenAPI specifications. Every service must support health check endpoints, structured JSON logging, and distributed tracing metadata headers (X-Request-ID).

Consequently, the Order Management service interacts with the ERP core database over port 5439. The database runs on PostgreSQL 15.4 to handle complex ACID transactions and structured JSON metadata queries.

Specifically, data pipelines ingest real-time sales transactions, process them via stream processors, and sync them to PostgreSQL and ChromaDB databases for reporting and analytics.

Additionally, data pipelines ingest real-time sales transactions, process them via stream processors, and sync them to PostgreSQL and ChromaDB databases for reporting and analytics.

Moreover, API schemas must follow OpenAPI specifications. Every service must support health check endpoints, structured JSON logging, and distributed tracing metadata headers (X-Request-ID).

In this context, vector embeddings are generated using local transformer models and stored in ChromaDB vector stores with cosine similarity configurations, allowing fast semantic similarity search.

Furthermore, the microservices architecture employs Kong API Gateway for routing client calls, securing endpoints via OAuth 2.0, and applying rate limits to protect backend services from denial of service.

Specifically, data pipelines ingest real-time sales transactions, process them via stream processors, and sync them to PostgreSQL and ChromaDB databases for reporting and analytics.

Furthermore, the Order Management service interacts with the ERP core database over port 5439. The database runs on PostgreSQL 15.4 to handle complex ACID transactions and structured JSON metadata queries.

In this context, the Order Management service interacts with the ERP core database over port 5439. The database runs on PostgreSQL 15.4 to handle complex ACID transactions and structured JSON metadata queries.

In this context, the Order Management service interacts with the ERP core database over port 5439. The database runs on PostgreSQL 15.4 to handle complex ACID transactions and structured JSON metadata queries.

Specifically, data pipelines ingest real-time sales transactions, process them via stream processors, and sync them to PostgreSQL and ChromaDB databases for reporting and analytics.

Consequently, the microservices architecture employs Kong API Gateway for routing client calls, securing endpoints via OAuth 2.0, and applying rate limits to protect backend services from denial of service.

In this context, the Order Management service interacts with the ERP core

database over port 5439. The database runs on PostgreSQL 15.4 to handle complex ACID transactions and structured JSON metadata queries.

Moreover, vector embeddings are generated using local transformer models and stored in ChromaDB vector stores with cosine similarity configurations, allowing fast semantic similarity search.

Moreover, the microservices architecture employs Kong API Gateway for routing client calls, securing endpoints via OAuth 2.0, and applying rate limits to protect backend services from denial of service.

Furthermore, data pipelines ingest real-time sales transactions, process them via stream processors, and sync them to PostgreSQL and ChromaDB databases for reporting and analytics.

Specifically, vector embeddings are generated using local transformer models and stored in ChromaDB vector stores with cosine similarity configurations, allowing fast semantic similarity search.

Consequently, API schemas must follow OpenAPI specifications. Every service must support health check endpoints, structured JSON logging, and distributed tracing metadata headers (X-Request-ID).

Specifically, the microservices architecture employs Kong API Gateway for routing client calls, securing endpoints via OAuth 2.0, and applying rate limits to protect backend services from denial of service.

In this context, API schemas must follow OpenAPI specifications. Every service must support health check endpoints, structured JSON logging, and distributed tracing metadata headers (X-Request-ID).

Additionally, vector embeddings are generated using local transformer models and stored in ChromaDB vector stores with cosine similarity configurations, allowing fast semantic similarity search.

Furthermore, the Order Management service interacts with the ERP core database over port 5439. The database runs on PostgreSQL 15.4 to handle complex ACID

transactions and structured JSON metadata queries.

In this context, the microservices architecture employs Kong API Gateway for routing client calls, securing endpoints via OAuth 2.0, and applying rate limits to protect backend services from denial of service.

Consequently, the Order Management service interacts with the ERP core database over port 5439. The database runs on PostgreSQL 15.4 to handle complex ACID transactions and structured JSON metadata queries.

In this context, data pipelines ingest real-time sales transactions, process them via stream processors, and sync them to PostgreSQL and ChromaDB databases for reporting and analytics.

Specifically, the microservices architecture employs Kong API Gateway for routing client calls, securing endpoints via OAuth 2.0, and applying rate limits to protect backend services from denial of service.

Furthermore, vector embeddings are generated using local transformer models and stored in ChromaDB vector stores with cosine similarity configurations, allowing fast semantic similarity search.

Additionally, vector embeddings are generated using local transformer models and stored in ChromaDB vector stores with cosine similarity configurations, allowing fast semantic similarity search.

In this context, the microservices architecture employs Kong API Gateway for routing client calls, securing endpoints via OAuth 2.0, and applying rate limits to protect backend services from denial of service.

Furthermore, API schemas must follow OpenAPI specifications. Every service must support health check endpoints, structured JSON logging, and distributed tracing metadata headers (X-Request-ID).

In this context, the microservices architecture employs Kong API Gateway for routing client calls, securing endpoints via OAuth 2.0, and applying rate limits to protect backend services from denial of service.

Additionally, the microservices architecture employs Kong API Gateway for routing client calls, securing endpoints via OAuth 2.0, and applying rate limits to protect backend services from denial of service.

Moreover, vector embeddings are generated using local transformer models and stored in ChromaDB vector stores with cosine similarity configurations, allowing fast semantic similarity search.

Specifically, data pipelines ingest real-time sales transactions, process them via stream processors, and sync them to PostgreSQL and ChromaDB databases for reporting and analytics.

Moreover, the microservices architecture employs Kong API Gateway for routing client calls, securing endpoints via OAuth 2.0, and applying rate limits to protect backend services from denial of service.

Consequently, the microservices architecture employs Kong API Gateway for routing client calls, securing endpoints via OAuth 2.0, and applying rate limits to protect backend services from denial of service.

Consequently, API schemas must follow OpenAPI specifications. Every service must support health check endpoints, structured JSON logging, and distributed tracing metadata headers (X-Request-ID).

Moreover, vector embeddings are generated using local transformer models and stored in ChromaDB vector stores with cosine similarity configurations, allowing fast semantic similarity search.

In this context, vector embeddings are generated using local transformer models and stored in ChromaDB vector stores with cosine similarity configurations, allowing fast semantic similarity search.

Furthermore, vector embeddings are generated using local transformer models and stored in ChromaDB vector stores with cosine similarity configurations, allowing fast semantic similarity search.

Consequently, data pipelines ingest real-time sales transactions, process them via stream processors, and sync them to PostgreSQL and ChromaDB databases for reporting and analytics.

Additionally, aPI schemas must follow OpenAPI specifications. Every service must support health check endpoints, structured JSON logging, and distributed tracing metadata headers (X-Request-ID).

Consequently, data pipelines ingest real-time sales transactions, process them via stream processors, and sync them to PostgreSQL and ChromaDB databases for reporting and analytics.

Specifically, the Order Management service interacts with the ERP core database over port 5439. The database runs on PostgreSQL 15.4 to handle complex ACID transactions and structured JSON metadata queries.

Consequently, aPI schemas must follow OpenAPI specifications. Every service must support health check endpoints, structured JSON logging, and distributed tracing metadata headers (X-Request-ID).

Furthermore, the Order Management service interacts with the ERP core database over port 5439. The database runs on PostgreSQL 15.4 to handle complex ACID transactions and structured JSON metadata queries.

Moreover, data pipelines ingest real-time sales transactions, process them via stream processors, and sync them to PostgreSQL and ChromaDB databases for reporting and analytics.

Furthermore, vector embeddings are generated using local transformer models and stored in ChromaDB vector stores with cosine similarity configurations, allowing fast semantic similarity search.

In this context, data pipelines ingest real-time sales transactions, process them via stream processors, and sync them to PostgreSQL and ChromaDB databases for reporting and analytics.

In this context, the microservices architecture employs Kong API Gateway for

routing client calls, securing endpoints via OAuth 2.0, and applying rate limits to protect backend services from denial of service.

Specifically, data pipelines ingest real-time sales transactions, process them via stream processors, and sync them to PostgreSQL and ChromaDB databases for reporting and analytics.

Moreover, the microservices architecture employs Kong API Gateway for routing client calls, securing endpoints via OAuth 2.0, and applying rate limits to protect backend services from denial of service.

Specifically, the microservices architecture employs Kong API Gateway for routing client calls, securing endpoints via OAuth 2.0, and applying rate limits to protect backend services from denial of service.

Consequently, API schemas must follow OpenAPI specifications. Every service must support health check endpoints, structured JSON logging, and distributed tracing metadata headers (X-Request-ID).

Moreover, data pipelines ingest real-time sales transactions, process them via stream processors, and sync them to PostgreSQL and ChromaDB databases for reporting and analytics.

Additionally, data pipelines ingest real-time sales transactions, process them via stream processors, and sync them to PostgreSQL and ChromaDB databases for reporting and analytics.

Additionally, data pipelines ingest real-time sales transactions, process them via stream processors, and sync them to PostgreSQL and ChromaDB databases for reporting and analytics.

Specifically, vector embeddings are generated using local transformer models and stored in ChromaDB vector stores with cosine similarity configurations, allowing fast semantic similarity search.

In this context, the microservices architecture employs Kong API Gateway for routing client calls, securing endpoints via OAuth 2.0, and applying rate limits to

protect backend services from denial of service.

Furthermore, the Order Management service interacts with the ERP core database over port 5439. The database runs on PostgreSQL 15.4 to handle complex ACID transactions and structured JSON metadata queries.

Additionally, aPI schemas must follow OpenAPI specifications. Every service must support health check endpoints, structured JSON logging, and distributed tracing metadata headers (X-Request-ID).

In this context, vector embeddings are generated using local transformer models and stored in ChromaDB vector stores with cosine similarity configurations, allowing fast semantic similarity search.

Furthermore, data pipelines ingest real-time sales transactions, process them via stream processors, and sync them to PostgreSQL and ChromaDB databases for reporting and analytics.

Furthermore, aPI schemas must follow OpenAPI specifications. Every service must support health check endpoints, structured JSON logging, and distributed tracing metadata headers (X-Request-ID).

In this context, vector embeddings are generated using local transformer models and stored in ChromaDB vector stores with cosine similarity configurations, allowing fast semantic similarity search.

Specifically, the Order Management service interacts with the ERP core database over port 5439. The database runs on PostgreSQL 15.4 to handle complex ACID transactions and structured JSON metadata queries.

In this context, data pipelines ingest real-time sales transactions, process them via stream processors, and sync them to PostgreSQL and ChromaDB databases for reporting and analytics.

Additionally, vector embeddings are generated using local transformer models and stored in ChromaDB vector stores with cosine similarity configurations, allowing fast semantic similarity search.

Additionally, vector embeddings are generated using local transformer models and stored in ChromaDB vector stores with cosine similarity configurations, allowing fast semantic similarity search.

Furthermore, aPI schemas must follow OpenAPI specifications. Every service must support health check endpoints, structured JSON logging, and distributed tracing metadata headers (X-Request-ID).

In this context, data pipelines ingest real-time sales transactions, process them via stream processors, and sync them to PostgreSQL and ChromaDB databases for reporting and analytics.

Furthermore, data pipelines ingest real-time sales transactions, process them via stream processors, and sync them to PostgreSQL and ChromaDB databases for reporting and analytics.

Furthermore, vector embeddings are generated using local transformer models and stored in ChromaDB vector stores with cosine similarity configurations, allowing fast semantic similarity search.

Moreover, the microservices architecture employs Kong API Gateway for routing client calls, securing endpoints via OAuth 2.0, and applying rate limits to protect backend services from denial of service.

In this context, data pipelines ingest real-time sales transactions, process them via stream processors, and sync them to PostgreSQL and ChromaDB databases for reporting and analytics.

Consequently, data pipelines ingest real-time sales transactions, process them via stream processors, and sync them to PostgreSQL and ChromaDB databases for reporting and analytics.

Moreover, aPI schemas must follow OpenAPI specifications. Every service must support health check endpoints, structured JSON logging, and distributed tracing metadata headers (X-Request-ID).

In this context, vector embeddings are generated using local transformer models and stored in ChromaDB vector stores with cosine similarity configurations, allowing fast semantic similarity search.

Consequently, the Order Management service interacts with the ERP core database over port 5439. The database runs on PostgreSQL 15.4 to handle complex ACID transactions and structured JSON metadata queries.

Specifically, vector embeddings are generated using local transformer models and stored in ChromaDB vector stores with cosine similarity configurations, allowing fast semantic similarity search.

Moreover, vector embeddings are generated using local transformer models and stored in ChromaDB vector stores with cosine similarity configurations, allowing fast semantic similarity search.

Consequently, data pipelines ingest real-time sales transactions, process them via stream processors, and sync them to PostgreSQL and ChromaDB databases for reporting and analytics.

Additionally, aAPI schemas must follow OpenAPI specifications. Every service must support health check endpoints, structured JSON logging, and distributed tracing metadata headers (X-Request-ID).

Moreover, aAPI schemas must follow OpenAPI specifications. Every service must support health check endpoints, structured JSON logging, and distributed tracing metadata headers (X-Request-ID).

Moreover, data pipelines ingest real-time sales transactions, process them via stream processors, and sync them to PostgreSQL and ChromaDB databases for reporting and analytics.

In this context, the microservices architecture employs Kong API Gateway for routing client calls, securing endpoints via OAuth 2.0, and applying rate limits to protect backend services from denial of service.

Moreover, the microservices architecture employs Kong API Gateway for routing

client calls, securing endpoints via OAuth 2.0, and applying rate limits to protect backend services from denial of service.

Additionally, data pipelines ingest real-time sales transactions, process them via stream processors, and sync them to PostgreSQL and ChromaDB databases for reporting and analytics.

In this context, the microservices architecture employs Kong API Gateway for routing client calls, securing endpoints via OAuth 2.0, and applying rate limits to protect backend services from denial of service.

Furthermore, data pipelines ingest real-time sales transactions, process them via stream processors, and sync them to PostgreSQL and ChromaDB databases for reporting and analytics.

Consequently, the Order Management service interacts with the ERP core database over port 5439. The database runs on PostgreSQL 15.4 to handle complex ACID transactions and structured JSON metadata queries.

Additionally, API schemas must follow OpenAPI specifications. Every service must support health check endpoints, structured JSON logging, and distributed tracing metadata headers (X-Request-ID).

Specifically, data pipelines ingest real-time sales transactions, process them via stream processors, and sync them to PostgreSQL and ChromaDB databases for reporting and analytics.

Specifically, the microservices architecture employs Kong API Gateway for routing client calls, securing endpoints via OAuth 2.0, and applying rate limits to protect backend services from denial of service.

Moreover, data pipelines ingest real-time sales transactions, process them via stream processors, and sync them to PostgreSQL and ChromaDB databases for reporting and analytics.

Moreover, data pipelines ingest real-time sales transactions, process them via stream processors, and sync them to PostgreSQL and ChromaDB databases for

reporting and analytics.

In this context, data pipelines ingest real-time sales transactions, process them via stream processors, and sync them to PostgreSQL and ChromaDB databases for reporting and analytics.

Furthermore, vector embeddings are generated using local transformer models and stored in ChromaDB vector stores with cosine similarity configurations, allowing fast semantic similarity search.

Moreover, the Order Management service interacts with the ERP core database over port 5439. The database runs on PostgreSQL 15.4 to handle complex ACID transactions and structured JSON metadata queries.

Consequently, data pipelines ingest real-time sales transactions, process them via stream processors, and sync them to PostgreSQL and ChromaDB databases for reporting and analytics.

Additionally, the microservices architecture employs Kong API Gateway for routing client calls, securing endpoints via OAuth 2.0, and applying rate limits to protect backend services from denial of service.

Moreover, API schemas must follow OpenAPI specifications. Every service must support health check endpoints, structured JSON logging, and distributed tracing metadata headers (X-Request-ID).

Specifically, vector embeddings are generated using local transformer models and stored in ChromaDB vector stores with cosine similarity configurations, allowing fast semantic similarity search.

Moreover, vector embeddings are generated using local transformer models and stored in ChromaDB vector stores with cosine similarity configurations, allowing fast semantic similarity search.

Additionally, data pipelines ingest real-time sales transactions, process them via stream processors, and sync them to PostgreSQL and ChromaDB databases for reporting and analytics.

Specifically, vector embeddings are generated using local transformer models and stored in ChromaDB vector stores with cosine similarity configurations, allowing fast semantic similarity search.

Consequently, the microservices architecture employs Kong API Gateway for routing client calls, securing endpoints via OAuth 2.0, and applying rate limits to protect backend services from denial of service.

Moreover, the microservices architecture employs Kong API Gateway for routing client calls, securing endpoints via OAuth 2.0, and applying rate limits to protect backend services from denial of service.

Additionally, the Order Management service interacts with the ERP core database over port 5439. The database runs on PostgreSQL 15.4 to handle complex ACID transactions and structured JSON metadata queries.

Furthermore, vector embeddings are generated using local transformer models and stored in ChromaDB vector stores with cosine similarity configurations, allowing fast semantic similarity search.

In this context, the microservices architecture employs Kong API Gateway for routing client calls, securing endpoints via OAuth 2.0, and applying rate limits to protect backend services from denial of service.

In this context, data pipelines ingest real-time sales transactions, process them via stream processors, and sync them to PostgreSQL and ChromaDB databases for reporting and analytics.

Moreover, data pipelines ingest real-time sales transactions, process them via stream processors, and sync them to PostgreSQL and ChromaDB databases for reporting and analytics.

Specifically, API schemas must follow OpenAPI specifications. Every service must support health check endpoints, structured JSON logging, and distributed tracing metadata headers (X-Request-ID).

Specifically, the Order Management service interacts with the ERP core database over port 5439. The database runs on PostgreSQL 15.4 to handle complex ACID transactions and structured JSON metadata queries.

Furthermore, the Order Management service interacts with the ERP core database over port 5439. The database runs on PostgreSQL 15.4 to handle complex ACID transactions and structured JSON metadata queries.

Additionally, data pipelines ingest real-time sales transactions, process them via stream processors, and sync them to PostgreSQL and ChromaDB databases for reporting and analytics.

Specifically, the Order Management service interacts with the ERP core database over port 5439. The database runs on PostgreSQL 15.4 to handle complex ACID transactions and structured JSON metadata queries.

Furthermore, vector embeddings are generated using local transformer models and stored in ChromaDB vector stores with cosine similarity configurations, allowing fast semantic similarity search.

Consequently, the microservices architecture employs Kong API Gateway for routing client calls, securing endpoints via OAuth 2.0, and applying rate limits to protect backend services from denial of service.

Moreover, the Order Management service interacts with the ERP core database over port 5439. The database runs on PostgreSQL 15.4 to handle complex ACID transactions and structured JSON metadata queries.

Moreover, the microservices architecture employs Kong API Gateway for routing client calls, securing endpoints via OAuth 2.0, and applying rate limits to protect backend services from denial of service.

Additionally, vector embeddings are generated using local transformer models and stored in ChromaDB vector stores with cosine similarity configurations, allowing fast semantic similarity search.

Furthermore, the Order Management service interacts with the ERP core database

over port 5439. The database runs on PostgreSQL 15.4 to handle complex ACID transactions and structured JSON metadata queries.

Specifically, API schemas must follow OpenAPI specifications. Every service must support health check endpoints, structured JSON logging, and distributed tracing metadata headers (X-Request-ID).

Specifically, API schemas must follow OpenAPI specifications. Every service must support health check endpoints, structured JSON logging, and distributed tracing metadata headers (X-Request-ID).

Moreover, the microservices architecture employs Kong API Gateway for routing client calls, securing endpoints via OAuth 2.0, and applying rate limits to protect backend services from denial of service.

Consequently, data pipelines ingest real-time sales transactions, process them via stream processors, and sync them to PostgreSQL and ChromaDB databases for reporting and analytics.

Consequently, data pipelines ingest real-time sales transactions, process them via stream processors, and sync them to PostgreSQL and ChromaDB databases for reporting and analytics.

Consequently, vector embeddings are generated using local transformer models and stored in ChromaDB vector stores with cosine similarity configurations, allowing fast semantic similarity search.

Furthermore, the microservices architecture employs Kong API Gateway for routing client calls, securing endpoints via OAuth 2.0, and applying rate limits to protect backend services from denial of service.

Specifically, vector embeddings are generated using local transformer models and stored in ChromaDB vector stores with cosine similarity configurations, allowing fast semantic similarity search.

Specifically, the microservices architecture employs Kong API Gateway for routing client calls, securing endpoints via OAuth 2.0, and applying rate limits to protect

backend services from denial of service.

Moreover, the microservices architecture employs Kong API Gateway for routing client calls, securing endpoints via OAuth 2.0, and applying rate limits to protect backend services from denial of service.

Furthermore, vector embeddings are generated using local transformer models and stored in ChromaDB vector stores with cosine similarity configurations, allowing fast semantic similarity search.

In this context, API schemas must follow OpenAPI specifications. Every service must support health check endpoints, structured JSON logging, and distributed tracing metadata headers (X-Request-ID).

In this context, the microservices architecture employs Kong API Gateway for routing client calls, securing endpoints via OAuth 2.0, and applying rate limits to protect backend services from denial of service.

Additionally, vector embeddings are generated using local transformer models and stored in ChromaDB vector stores with cosine similarity configurations, allowing fast semantic similarity search.

Consequently, the Order Management service interacts with the ERP core database over port 5439. The database runs on PostgreSQL 15.4 to handle complex ACID transactions and structured JSON metadata queries.

In this context, the Order Management service interacts with the ERP core database over port 5439. The database runs on PostgreSQL 15.4 to handle complex ACID transactions and structured JSON metadata queries.

Moreover, the microservices architecture employs Kong API Gateway for routing client calls, securing endpoints via OAuth 2.0, and applying rate limits to protect backend services from denial of service.

Moreover, the Order Management service interacts with the ERP core database over port 5439. The database runs on PostgreSQL 15.4 to handle complex ACID transactions and structured JSON metadata queries.

Specifically, data pipelines ingest real-time sales transactions, process them via stream processors, and sync them to PostgreSQL and ChromaDB databases for reporting and analytics.

Additionally, API schemas must follow OpenAPI specifications. Every service must support health check endpoints, structured JSON logging, and distributed tracing metadata headers (X-Request-ID).

Furthermore, data pipelines ingest real-time sales transactions, process them via stream processors, and sync them to PostgreSQL and ChromaDB databases for reporting and analytics.

Additionally, vector embeddings are generated using local transformer models and stored in ChromaDB vector stores with cosine similarity configurations, allowing fast semantic similarity search.

Additionally, the Order Management service interacts with the ERP core database over port 5439. The database runs on PostgreSQL 15.4 to handle complex ACID transactions and structured JSON metadata queries.

Additionally, API schemas must follow OpenAPI specifications. Every service must support health check endpoints, structured JSON logging, and distributed tracing metadata headers (X-Request-ID).

Additionally, the microservices architecture employs Kong API Gateway for routing client calls, securing endpoints via OAuth 2.0, and applying rate limits to protect backend services from denial of service.

Additionally, vector embeddings are generated using local transformer models and stored in ChromaDB vector stores with cosine similarity configurations, allowing fast semantic similarity search.

Moreover, vector embeddings are generated using local transformer models and stored in ChromaDB vector stores with cosine similarity configurations, allowing fast semantic similarity search.

Furthermore, the microservices architecture employs Kong API Gateway for routing client calls, securing endpoints via OAuth 2.0, and applying rate limits to protect backend services from denial of service.

Furthermore, the Order Management service interacts with the ERP core database over port 5439. The database runs on PostgreSQL 15.4 to handle complex ACID transactions and structured JSON metadata queries.

Specifically, the microservices architecture employs Kong API Gateway for routing client calls, securing endpoints via OAuth 2.0, and applying rate limits to protect backend services from denial of service.

Consequently, the microservices architecture employs Kong API Gateway for routing client calls, securing endpoints via OAuth 2.0, and applying rate limits to protect backend services from denial of service.

Additionally, vector embeddings are generated using local transformer models and stored in ChromaDB vector stores with cosine similarity configurations, allowing fast semantic similarity search.

Specifically, the microservices architecture employs Kong API Gateway for routing client calls, securing endpoints via OAuth 2.0, and applying rate limits to protect backend services from denial of service.

Moreover, API schemas must follow OpenAPI specifications. Every service must support health check endpoints, structured JSON logging, and distributed tracing metadata headers (X-Request-ID).

Furthermore, API schemas must follow OpenAPI specifications. Every service must support health check endpoints, structured JSON logging, and distributed tracing metadata headers (X-Request-ID).

Additionally, the microservices architecture employs Kong API Gateway for routing client calls, securing endpoints via OAuth 2.0, and applying rate limits to protect backend services from denial of service.

Additionally, data pipelines ingest real-time sales transactions, process them via

stream processors, and sync them to PostgreSQL and ChromaDB databases for reporting and analytics.

Specifically, the microservices architecture employs Kong API Gateway for routing client calls, securing endpoints via OAuth 2.0, and applying rate limits to protect backend services from denial of service.

Furthermore, API schemas must follow OpenAPI specifications. Every service must support health check endpoints, structured JSON logging, and distributed tracing metadata headers (X-Request-ID).

Specifically, vector embeddings are generated using local transformer models and stored in ChromaDB vector stores with cosine similarity configurations, allowing fast semantic similarity search.

Moreover, API schemas must follow OpenAPI specifications. Every service must support health check endpoints, structured JSON logging, and distributed tracing metadata headers (X-Request-ID).

Furthermore, data pipelines ingest real-time sales transactions, process them via stream processors, and sync them to PostgreSQL and ChromaDB databases for reporting and analytics.

In this context, the Order Management service interacts with the ERP core database over port 5439. The database runs on PostgreSQL 15.4 to handle complex ACID transactions and structured JSON metadata queries.

Consequently, the microservices architecture employs Kong API Gateway for routing client calls, securing endpoints via OAuth 2.0, and applying rate limits to protect backend services from denial of service.

Furthermore, vector embeddings are generated using local transformer models and stored in ChromaDB vector stores with cosine similarity configurations, allowing fast semantic similarity search.

Furthermore, the Order Management service interacts with the ERP core database over port 5439. The database runs on PostgreSQL 15.4 to handle complex ACID

transactions and structured JSON metadata queries.

Furthermore, data pipelines ingest real-time sales transactions, process them via stream processors, and sync them to PostgreSQL and ChromaDB databases for reporting and analytics.

Consequently, the Order Management service interacts with the ERP core database over port 5439. The database runs on PostgreSQL 15.4 to handle complex ACID transactions and structured JSON metadata queries.

In this context, the microservices architecture employs Kong API Gateway for routing client calls, securing endpoints via OAuth 2.0, and applying rate limits to protect backend services from denial of service.

Additionally, API schemas must follow OpenAPI specifications. Every service must support health check endpoints, structured JSON logging, and distributed tracing metadata headers (X-Request-ID).

Consequently, the microservices architecture employs Kong API Gateway for routing client calls, securing endpoints via OAuth 2.0, and applying rate limits to protect backend services from denial of service.

Specifically, vector embeddings are generated using local transformer models and stored in ChromaDB vector stores with cosine similarity configurations, allowing fast semantic similarity search.

In this context, the Order Management service interacts with the ERP core database over port 5439. The database runs on PostgreSQL 15.4 to handle complex ACID transactions and structured JSON metadata queries.

Specifically, data pipelines ingest real-time sales transactions, process them via stream processors, and sync them to PostgreSQL and ChromaDB databases for reporting and analytics.

In this context, the Order Management service interacts with the ERP core database over port 5439. The database runs on PostgreSQL 15.4 to handle complex ACID transactions and structured JSON metadata queries.

Moreover, the microservices architecture employs Kong API Gateway for routing client calls, securing endpoints via OAuth 2.0, and applying rate limits to protect backend services from denial of service.

Consequently, vector embeddings are generated using local transformer models and stored in ChromaDB vector stores with cosine similarity configurations, allowing fast semantic similarity search.

In this context, the Order Management service interacts with the ERP core database over port 5439. The database runs on PostgreSQL 15.4 to handle complex ACID transactions and structured JSON metadata queries.

Specifically, data pipelines ingest real-time sales transactions, process them via stream processors, and sync them to PostgreSQL and ChromaDB databases for reporting and analytics.

Moreover, vector embeddings are generated using local transformer models and stored in ChromaDB vector stores with cosine similarity configurations, allowing fast semantic similarity search.

Consequently, data pipelines ingest real-time sales transactions, process them via stream processors, and sync them to PostgreSQL and ChromaDB databases for reporting and analytics.

Additionally, the microservices architecture employs Kong API Gateway for routing client calls, securing endpoints via OAuth 2.0, and applying rate limits to protect backend services from denial of service.